

CS 7394

```
In [1]: import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
```

Source: World Bank (<https://data.worldbank.org/indicator/SP.POP.TOTL>)

```
Out[2]:
```

	Country Name	1960	1961	1962	1963	1964	1965	1966	
0	Afghanistan	8996967.0	9169406.0	9351442.0	9543200.0	9744772.0	9956318.0	10174840.0	1039
1	Albania	1608800.0	1659800.0	1711319.0	1762621.0	1814135.0	1864791.0	1914573.0	196
2	Algeria	11057864.0	11336336.0	11619828.0	11912800.0	12221675.0	12550880.0	12902626.0	1327
3	American Samoa	20127.0	20605.0	21246.0	22029.0	22850.0	23675.0	24473.0	2
4	Andorra	13410.0	14378.0	15379.0	16407.0	17466.0	18542.0	19646.0	2

5 rows × 62 columns

Out[3]: (217, 62)

Country Name	1960	1961	1962	1963	1964	1965	1966	1967
61 Eritrea	1007586.0	1033320.0	1060489.0	1088859.0	1118152.0	1148188.0	1178875.0	1210304.0
106 Kuwait	269026.0	300581.0	337346.0	378756.0	423900.0	472032.0	523169.0	577164.0
213 West Bank and Gaza	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

3 rows × 62 columns

```
In [5]: data.dropna(inplace=True)
```

```
In [6]: data.shape
```

```
Out[6]: (214, 62)
```

```
In [7]: data = data.drop(columns=['Country Name'])
```

Model Fit and Predict

```
In [8]: def build_model(train, modType):
        x = train.iloc[:, 0].values.reshape(-1,1)
        y = train.iloc[:, 1].values.reshape(-1,1)
        model = modType().fit(x,y)
        return model

        def predict(model, year):
            return model.predict([[year]])[0][0]
```

Predict World Population in 2122

Method #1:

Create one model for the total population

```
In [9]: data2 = pd.read_csv('totalPop.csv')
        data2.head()
```

```
Out[9]:
```

	Series Name	Country Name	1960	1961	1962	1963	1964	1965
0	Population, total	World	3032156070	3071596055	3124561005	3189655687	3255145692	3322046795

1 rows × 63 columns

```
In [10]: data2 = data2.drop(columns=['Series Name', 'Country Name'])
```

```
In [11]: data_tuples = list(zip(data2.columns.values, data2.iloc[0, :]))
        tempDf = pd.DataFrame(data_tuples, columns=['Year', 'Population'])
```

```
In [12]: linReg = build_model(tempDf, LinearRegression)

predLin1 = predict(linReg, 2122)
print("Linear Regression #1 population prediction = ", np.floor(predLin1))
```

```
Linear Regression #1 population prediction = 15953694673.0
```

Model plot

```
In [13]: from matplotlib import pyplot as plt
import seaborn as sns

X = tempDf.iloc[:, 0].values
Y = tempDf.iloc[:, 1].values
x = X.reshape(-1,1)
```

```
In [14]: y_pred = linReg.predict(x)
y_pred = y_pred.reshape(1,-1)
```

```
In [15]: import matplotlib.ticker as ticker

plt.figure(figsize=(15, 10))

sns.set()


ax = sns.scatterplot(X, Y, color = 'red')
sns.lineplot(X, y = y_pred[0])
ax.xaxis.set_major_locator(ticker.MultipleLocator(5))
ax.xaxis.set_major_formatter(ticker.ScalarFormatter(-1960,-2020))
plt.title('Population Prediction', fontsize=20)
plt.xlabel('Years', fontsize=20)
plt.ylabel('Population', fontsize=20)
plt.legend(['Prediction', 'Population'], fontsize=15)
plt.savefig("model.png")

plt.show()
```

/home/sam/anaconda3/envs/aplML/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(
/home/sam/anaconda3/envs/aplML/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

cit keyword will result in an error or misinterpretation.
warnings.warn(



```
In [16]: from sklearn.metrics import mean_squared_error

rms = mean_squared_error(Y, linReg.predict(x), squared=False)
rms
print("RMSE of the linear regression model =", rms)
```

RMSE of the linear regression model = 41276568.06427081

Method #2:

- Model the population growth of each country
- Predict the population of one country
- Sum the predictedtions

```
In [17]: predLin2 = 0
for index in range(data.shape[0]):
    data_tuples = list(zip(data.columns.values, data.iloc[index, :]))
    data_tuples

    tempDf = pd.DataFrame(data_tuples, columns=['Year', 'Population'])
    model = build_model(tempDf, LinearRegression)
```

```

pred = predict(model, 2122)
predLin2 = predLin2 + pred

print("Linear Regression #1 population prediction = ", np.floor(predLin2))

```

Linear Regression #1 population prediction = 15874354113.0

Method #2 yields a slightly lower prediction than Method #1 which makes it the better method, but still not enough.

Ridge Regression

```

In [18]: from sklearn.linear_model import Ridge
         from sklearn.linear_model import Lasso

```

Method #1:

```

In [19]: data_tuples = list(zip(data2.columns.values, data2.iloc[0, : ]))
         tempDf = pd.DataFrame(data_tuples, columns=['Year', 'Population'])

         x = tempDf.iloc[:, 0].values.reshape(-1,1)
         y = tempDf.iloc[:, 1].values.reshape(-1,1)

         ridge = Ridge(alpha=1.0)
         # fit
         ridge.fit(x, y)
         # predict

         predRidge1 = predict(ridge, 2122)

         predRidge1

```

Out[19]: 15953131404.865936

```

In [20]: rms2 = mean_squared_error(Y, ridge.predict(x), squared=False)
         print("RMSE of the ridge regression model =", rms2)

```

RMSE of the ridge regression model = 41276636.44146911

Method #2:

```

In [21]: predRidge2 = 0
         for index in range(data.shape[0]):
             data_tuples = list(zip(data.columns.values, data.iloc[index, : ]))
             data_tuples

```

```
tempDf = pd.DataFrame(data_tuples, columns=['Year', 'Population'])
model = build_model(tempDf, Ridge)

pred = predict(model, 2122)
predRidge2 = predRidge2 + pred
#print(pred)
print("Ridge Regression population prediction = ", predRidge2)
```

```
Ridge Regression population prediction = 15873793727.353664
```

I was curious to see how differently ridge regression will perform on this problem. In both methods it produced almost the same result as the linear regression models, with prediction values lower than that of LR by around 500K.

Conclusion

Linear Regression is not the optimal solution for our problem. It tends to provide high predictions because it follows population growth data from over the past 80 years only, during which world population boomed. Between 1960 and 2020 (the dataset limits), the world population more than doubled, growing from 3 billion to almost 8 billion. It is reasonable for a linear model to predict that the population doubles in 100 years, given only these information.

The information extracted from the population data is not enough to make a reliable prediction 100 years in the future. We probably need to look for more data features that contribute to population growth/decline. Such features include (but not limited to): fertility rates, mortality rates, death rates, gross reproduction rates, migration rates, deforestation, climate change, etc.... These features could be consolidated with dimensionality reduction algorithms to extract the most information out of the data that indicate a trend. We then train a neural network on the extracted features. This approach would provide more reliable and data driven predictions.